

STAT1301 - Lecture Notes

Ismael Khan

Chapter 1

Statistics Stuff

1.1 Lecture 5

$$\bar{x} = \frac{x_1 + x_2 + x_3 + \dots + x_n}{n} = \frac{\sum_{i=1}^n x_i}{n}$$

The value $b = \bar{x}$ minimizes the squared deviations from the data:

$$\sum_{i=1}^n (x_i - b)^2$$

Prove this Observe that

$$\begin{aligned} g(b) &= (170 - b)^2 + (182 - b)^2 + (160 - b)^2 \\ g(b) &= 3b^2 - \beta b - \gamma b - \end{aligned}$$

We can take the derivative to find the minimum

$$\begin{aligned} g(b) &= \sum_{i=1}^n (x_i - b)^2 \\ g'(b) &= \sum_{i=1}^n 2(x_i - b) \cdot -1 = 0 \\ &= test \end{aligned}$$

1.1.1 Robustness

What would happen if the third student had given their value in metres? Outliers can have a big effect on the sample mean. The median is also more informative for skewed data. Thus visualising is important before using means.

1.1.2 Sample Variance

Consider the three height values again: 170cm, 185cm, 162cm. How do we measure the spread of these values?

One way to measure spread is via the sample variance:

$$\beta = \frac{\sum_{i=1}^n (x_i - \bar{x})^2}{n - 1}$$

1.1.3 Sample standard deviation

Note that the sample variance is on the squared scale of the original observations. To get a measure of spread on the original scale, we use the function

$$\sigma = \sqrt{\frac{\sum_{i=1}^n (x_i - \bar{x})^2}{n - 1}}$$

Take an average height of people to be 170cm. A person of height 182cm leaves the room. The mean height will decrease. We can show this.

$$\frac{x_1 + x_2 + x_3 + \dots + x_n}{n} = 170$$

Without loss of generality, that $x_1 = 180$, then

$$\frac{x_2 + \dots + x_n}{n} = \frac{x_1 + x_2 + \dots + x_n - x_1}{n - 1}$$

$$\frac{n \times 170 - 180}{n - 1} < \frac{n \times 170 - 170}{n - 1} = \frac{170(n - 1)}{n - 1} = 170$$

1.1.4 Uniform random numbers

1.2 Lecture 6

1.2.1 Modelling Relationships

We often model relationships as a trend in the mean response plus variability about that trend.

$$\gamma(x) = g(x) + \varepsilon$$

For γ being the response, x being the explanation variable and ε being the random error

1.2.2 Describing Relations

How much would describe the relationship between the father and son?

1.2.3 Pearson Correlation

The Pearson correlation coefficient r , measures the strength of a linear relationship. If the points in the scatter plot are (x_1, y_1) , (x_2, y_2) , (x_n, y_n) then

$$r = \frac{1}{n - 1} \sum_{i=1}^n \left(\frac{x_i - \bar{x}}{s_x} \right) \left(\frac{y_i - \bar{y}}{s_y} \right)$$

It is important to remember that r is only appropriate for linear relationships between variables.

1.2.4 Least-Squared Lines

Once we have determined that a straight line relationship may be appropriate we can go ahead and fit a line of best fit to the data. For any given line $b_0 + b_1x$, how to judge how well it fits the data? We can look at the sum of the squared prediction error (or sum of squared deviations)

$$\sum_{i=1}^n (y_i - [b_0 + b_1x_i])^2$$

The line $b_0 + b_1x_1$, that minimises this is called the least-squares line.

1.3 Lecture 7

Probably skipped this to eat food

1.4 Lecture 8

I don't know anymore. What is the purpose of life. Who even is Joe Mama.

1.4.1 Simple Probability Models

The simplest probability model is when the sample spaces is a finite collection of outcomes, all being equally likely.

Example 1 (Rolling a Die). Consider rolling a fair 6 sided die. What is Ω ?

$$\Omega = \{1, 2, 3, 4, 5, 6\}$$

Then what is the probability \mathbb{P} . For a given $A \subset \Omega$, we assign the probability $\mathbb{P}(A)$ as

$$\mathbb{P}(A) = \frac{|A|}{|\Omega|}, \quad A \subset \Omega$$

Thus for $|\Omega| = 6$,

$$\mathbb{P}(A) = \frac{|A|}{6}, \quad A \subset \Omega$$

Drawing balls from an urn. The Urn can be counted in 4 possible urn experiments

1. Ordered, with replacement (Notes missing)
2. Ordered, without replacement
3. Unordered, without replacement

Consider a horse race with 8 horses. How many ways are there to gamble on placings (1st,2nd,3rd).

4. Unordered, with replacement

Theorem 1. Let $X \sim U[a, b]$. Then,

$$\begin{aligned} 1. \mathbb{E}(X) &= \frac{(a+b)}{2} \\ 2. \text{Var}(X) &= \frac{(b-a)^2}{12} \end{aligned}$$

1.5 Lecture 9 - Conditional Probability and Independence

1.5.1 Conditional Probability

How do we update the probability of an event A when we know that some other event B has occurred? If we are given that B has occurred, then A will occur if and only if $A \cap B$ occurs.

The relative chance of A occurring given B has occurred is therefore $\frac{\mathbb{P}(A \cap B)}{\mathbb{P}(B)}$. Which is the *conditional probability* of A given B .

Definition 1. The **conditional probability** of A given B is defined as.

$$\mathbb{P}(A | B) = \frac{\mathbb{P}(A \cap B)}{\mathbb{P}(B)}$$

Example 2. Suppose we have rolled a fair six-sided die. What is the conditional probability that we get a "4" given we know that we rolled an *even number*.

Intuition implies that the probability is $\frac{1}{3}$. We can formally show this however.

Let B be getting an even number $= \{2, 4, 6\}$ and A to be getting a 4 $= \{4\}$. We already know that $\mathbb{P}(B) = \frac{3}{6}$. Now $A \cap B = \{4\}$, so that $\mathbb{P}(A \cap B) = \frac{1}{6}$. Thus,

$$\mathbb{P}(A | B) = \frac{\mathbb{P}(A \cap B)}{\mathbb{P}(B)} = \frac{\frac{1}{6}}{\frac{3}{6}} = \frac{1}{3}$$

1.5.2 Product Rule

Rearranging the definition of the conditional probability gives us an expression for the product rule.

$$\mathbb{P}(A \cap B) = \mathbb{P}(B)\mathbb{P}(A | B) = \mathbb{P}(A)\mathbb{P}(B | A)$$

More generally,

$$\mathbb{P}(A_1 \cap \dots \cap A_n) = \mathbb{P}(A_1)\mathbb{P}(A_2 | A_1)\mathbb{P}(A_3 | A_1 \cap A_2) \dots \mathbb{P}(A_n | A_1 \cap \dots \cap A_{n-1})$$

The crucial usage of this is that in many cases conditional probabilities are easy to figure out.

Example 3. We draw consecutively 3 balls from an urn with 5 white and 5 black balls, without putting them back. What is the probability that all drawn balls will be black? Let A_i be the event that the i -th ball is black. We wish to find the probability of $A_1 \cap A_2 \cap A_3$ (also written as $A_1 A_2 A_3$), which by the product rule is

$$\mathbb{P}(A_1)\mathbb{P}(A_2 | A_1)\mathbb{P}(A_3 | A_1 \cap A_2) = \frac{5}{10} \frac{4}{9} \frac{3}{8} \approx 0.083$$

1.5.3 Independence

If the joint probability of two events A and B happens to factorise into the product of the two individual probabilities, i.e.

$$\mathbb{P}(A \cap B) = \mathbb{P}(A)\mathbb{P}(B)$$

Then we say that the events A and B are independent. In other words, knowledge of A gives no additional knowledge to B and vice versa. This implies that

$$\mathbb{P}(A | B) = \mathbb{P}(A)$$

and

$$\mathbb{P}(B | A) = \mathbb{P}(B)$$

Example 4. Example 3 is used however 3 balls are chosen **with** putting them back. So by independence,

$$\mathbb{P}(A_1)\mathbb{P}(A_2 | A_1)\mathbb{P}(A_3 | A_1 \cap A_2) = \frac{5}{10} \frac{5}{10} \frac{5}{10} \approx 0.125$$

1.6 Lecture 10 - Random Variables

A random variable can be viewed as measurement of a random experiment that becomes available **tomorrow**. However all the thinking about the measurements can be carried out **today**. We denote random variables with capital letters X, X_1, X_2, Y, Z etc.

Example 5. Some random variables may be

1. The number of defective transistors out of 100 inspected ones.
2. The number of bugs in a computer program
3. The amount of rain in a certain location in June.
4. The amount of time needed for an operation.

We distinguish between discrete and continuous random variables:

- **Discrete** random variables can only take *countably many* values.
- **Continuous** random variables can take a *continuous range* of values, for example, any value on the positive real line \mathbb{R}_+ .

1.6.1 Probability distribution

Let X be a random variable. We would like to designate the probabilities of events such as $\{X = x\}$ and $\{a \leq X \leq b\}$. If we can specify all probabilities involving X , we say that we have determined the probability distribution of X .

A way to specify the P.D is to give all probabilities of the form $\{X \leq x\}$.

Definition 2. The **cumulative distribution function** (cdf) of a random variable X is the function F defined by

$$F(x) = \mathbb{P}(X \leq x), x \in \mathbb{R}$$

1.6.2 Cumulative distribution function (cdf)

Note that any cdf is increasing (if $x \leq y$ then $F(x) \leq F(y)$) and lies inbetween 0 and 1. We can use any function F with these properties to specify the distrubution of a random variable X

1.6.3 Probability mass function

Definition 3. A random variable X is said to have a discrete distrubution if $\mathbb{P}(X = x_i) > 0$, $i = 1, 2, \dots$ for some finite or countable set of values x_1, x_2, \dots such that $\sum_i \mathbb{P}(X = x_i) = 1$. The **probability mass function (pmf)** of X is the function f defined by $f(x) = \mathbb{P}(X = x)$.

By the sum rule, if we know $f(x) \forall x$ then we can calculate all possible probabilties involving X

$$\mathbb{P}(X \in B) = \sum_{x \in B} f(x)$$

Note. $\{X \in B\}$ should be read as X is an elemnt of B

Definition 4. A random variable X with cdf F is said to have a continuous distrubution if there exists a positive function f with total integral I such that $\forall a < b$

$$\mathbb{P}(a < X \leq b) = F(b) - F(a) = \int_a^b f(u) du$$

Function f is called the **probability distribution function (pdf)** of X

1.6.4 Calculating probabilities

Once we know the pdf, we can calculate any probability taht X lies in some set B by means of integration

$$\mathbb{P}(X \in B) = \int_B f(x) dx$$

1.6.5 Relationship between cdf and pdf

Suppose f and F are the pdf and cdf of a continuous random variable X respectively. Then F is an anti-derivative of f

$$F(x) = \mathbb{P}(X \leq x) = \int_{-\infty}^x f(u) du$$

Conversely, f is the derivative of the cdf F

$$f(x) = \frac{d}{dx} F(x) = F'(x)$$

1.6.6 Density

In the continuous case, $f(x) \neq \mathbb{P}(X = x)$ because the latter is 0 $\forall x$. Instead we intepret $f(x)$ as the density of the probability at distrubution x , in the sense that for any small h .

$$\mathbb{P}(x \leq X \leq x+h) = \int_x^{x+h} f(u) du \approx hf(x)$$

Note. $\mathbb{P}(x \leq X \leq x+h) = \mathbb{P}(x < X \leq x+h)$ in this case.

Example 6. Draw a random number X from the interval of real numbers $[0, 2]$, where each number is equally likely to be drawn. What are the pdf and cdf F of X

We have

$$\mathbb{P}(X \leq x) = F(x) = \begin{cases} 0 & \text{if } x < 0 \\ \frac{x}{2} & \text{if } 0 \leq x \leq 2 \\ 1 & \text{if } x > 2 \end{cases}$$

By differentiating F we find

$$f(x) = \begin{cases} \frac{1}{2} & \text{if } 0 \leq x \leq 2 \\ 0 & \text{otherwise} \end{cases}$$

Note that this density is constant on the interval $[0, 2]$ and zero elsewhere. Reflecting the fact that each point in $[0, 2]$ are equally **equally likely to be drawn**.

1.6.7 Expectation (discrete case)

Although all probability info about a random variable is contained in its cdf or pmf/pdf, it is often useful to consider various numerical characteristics of a random variable.

Definition 5. Let X be a *discrete* random variable with pmf f . The **expectation** (or expected value) of X , denoted as $\mathbb{E}(X)$ is defined as

$$\mathbb{E}(X) = \sum_x x\mathbb{P}(X = x) = \sum_x xf(x)$$

The expectation is thus a "weighted average" of the values that X can take.

1.6.8 Expectation (continuous case)

Definition 6. Let X be a *continuous* random variable with pdf f , The **expectation** (or expected value) of X , denoted $\mathbb{E}X$, is defined as

$$\mathbb{E}(X) = \int_{-\infty}^{\infty} xf(x) dx$$

1.6.9 Expectation of a function of a random variable

If X is a random variable, then a function of X such as X^2 or $\sin(X)$ is also a random variable.

Theorem 2 (Expectation of a Function of a Random Variable). *If X is discrete with pdf f , then for any real values function g*

$$\mathbb{E}(g(X)) = \sum_x g(x)f(x)$$

Replace the integral with a sum for the discrete case

$$\mathbb{E}(g(X)) = \int_{-\infty}^{\infty} g(x)f(x) dx$$

1.6.10 Linearity of expectation

An important consequence of Theorem 2 is that the expectation is "linear".

Theorem 3. *For any real numbers a and b , and functions g and h .*

1. $\mathbb{E}(aX + b) = a\mathbb{E}(X) + b$
2. $\mathbb{E}(g(X) + h(X)) = \mathbb{E}(g(X)) + \mathbb{E}(h(X))$

Proof of Theorem 3. We show it for the discrete case. Suppose X has pmf f . The first statement follows from

$$\mathbb{E}(aX + b) = \sum_x (ax + b)f(x) = a \sum_x xf(x) + b \sum_x f(x) = a\mathbb{E}(X) + b$$

Similarly, the second statement follows from

$$\begin{aligned}\mathbb{E}(g(X) + h(X)) &= \sum_x (g(x) + h(x))f(x) = \sum_x h(x)f(x) \\ &\quad + \sum_x g(x)f(x) = \mathbb{E}(h(X)) + \mathbb{E}(g(X))\end{aligned}$$

□

1.6.11 Variance

Another useful characteristic of the distribution of X is the variance of X . This number is sometimes denoted as σ^2 and measures the *spread* of the distribution of X .

Definition 7. The **variance** of a random variable X , denoted $\text{Var}(X)$, is defined as

$$\text{Var}(X) = \mathbb{E}(X - \mu)^2$$

where $\mu = \mathbb{E}(X)$. The square root of the variance is called the **standard deviation**. The number $\mathbb{E}X^r$ is called the r -th **moment** of X .

1.6.12 Properties of Variance

Theorem 4. For any random variable X , the following properties hold for variance

1. $\text{Var}(X) = \mathbb{E}X^2 - (\mathbb{E}X)^2$
2. $\text{Var}(a + bX) = b^2 \text{Var}(X)$

1.6.13 Valuable Conclusions

- The probability distribution of a random variable is completely specified by its cdf (cumulative distribution function). For discrete random variables, it is more useful to specify distribution with pmf (probability mass function). For continuous random variables, use pdf (probability density function)
- The expectation (expected value) of a random variable is the weighted average of the values that random variable can take. It measures the locality of the distribution of a random variable.
- The variance is the expected squared distance from the random variable to its expected value. As a consequence, it is a measure of the spread of distribution of a random variable.

Note. The lecture 11 (scheduled on a Wednesday) was cancelled due to the Ekka holiday. So all lectures proceeding from this point are one less than their *true* lecture number.

1.7 Lecture 11 - Bernoulli and binomial distribution

1.7.1 Discrete random variables

Discrete random variables have only a finite or countable number of different values. Probabilities are calculated by summing the probability mass function (pmf). The most frequently used discrete probability distributions are the




Figure 1.1: Probability mass function (pmf)

Bernoulli and **Binomial** distributions.

1.7.2 Bernoulli Trial

A **Bernoulli Trial** is a random experiment that has only two outcomes (success - 1, failure - 0). The random variable X is called the Bernoulli variable.

An example may be a coin, two possible outcomes Heads or Tails, can be represented can be modelled as a Bernoulli variable attributing 1 for heads and 0 for tails.

Definition 8. A random variable X is said to have a **Bernoulli** distribution with success probability p if X can only assume the values 1 and 0 with probabilities

$$\mathbb{P}(X = 1) = p \text{ and } \mathbb{P}(X = 0) = 1 - p$$

We write $X \sim \text{Ber}(p)$

Theorem 5. Let $X \sim \text{Ber}(p)$. Then,

1. $\mathbb{E}(X) = p$
2. $\text{Var}(X) = p(1 - p)$

Lets prove this.

Proof. 1.

$$\begin{aligned}\mathbb{E}(X) &= \sum_x x f(x) \\ &= 0 \cdot (1 - p) + 1 \cdot p = p\end{aligned}$$

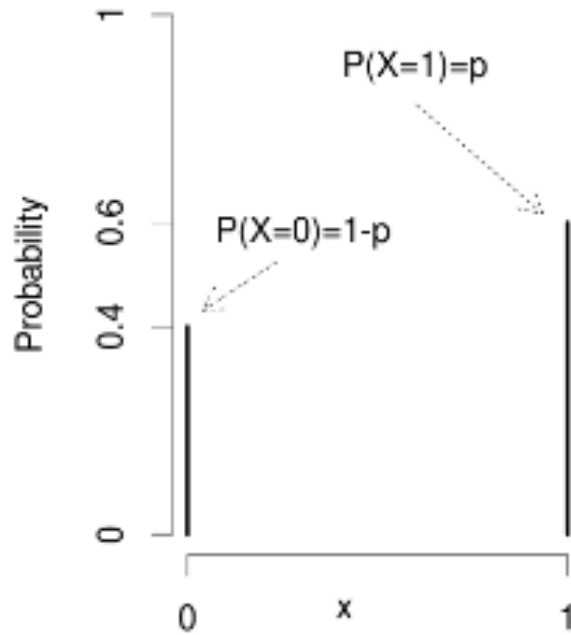


Figure 1.2: Probability mass function for the Bernoulli distribution with parameter p (the case $p = 0.6$ is shown)

2.

$$\begin{aligned}
 \text{Var}(X) &= \mathbb{E}(X - \mathbb{E}X)^2 \\
 &= \mathbb{E}(X - p)^2 \\
 &= \mathbb{E}(X^2 - 2pX + p^2) \\
 &= \mathbb{E}X^2 + (\mathbb{E}(-2pX)) + \mathbb{E}p^2 \\
 &= \mathbb{E}X^2 - 2p\mathbb{E}X + p^2 \\
 &= \mathbb{E}X^2 - p^2
 \end{aligned}$$

Note that $\mathbb{E}X^2 = 0^2(1 - p) + 1^2 \cdot p = p$

$$= p - p^2 = p(1 - p)$$

□

1.7.3 Counting possible outcomes

If we observe the complete sequence of 100 tosses, there are 2^{100} possible outcomes of the experiment, which are all equally likely (given a fair coin).

To calculate the probability of having exactly x successes we need to see how many of the possible outcomes have exactly x ones and $100 - x$ zeros.

There are $\binom{100}{x}$ of these, because we have to choose exactly x positions for the 1s out of 100 possible positions. We derive,

$$\mathbb{P}(X = x) = \frac{\binom{100}{x}}{2^{100}}, x = 0, 1, 2, \dots, 100$$

This is an example of a Binomial distribution. We can now calculate probabilities of interest such as $\mathbb{P}(X \leq 60)$. We

need to evaluate

$$\mathbb{P}(X \leq 60) = \sum_{x=60}^{100} \frac{\binom{100}{x}}{2^{100}} = \frac{\sum_{x=60}^{100} \binom{100}{x}}{2^{100}}$$

Which can be done in R with one line:

```
> sum(choose)(100,60:100))/2^(100)
```

```
[1] 0.02844397
```

1.7.4 Number of successes with a general coin

Since the coin tosses are independent, we can use the product rule to find that the probability of having a sequence of x heads and $n - x$ tails is $p^x(1 - p)^{n-x}$. However as there are still $\binom{n}{x}$ of these sequences. We can formalise this:

Definition 9. A random variable X is said to have a **binomial** distribution with parameters n and p if X can only assume the integer values of $x = 0, 1, \dots, n$, with probabilities

$$P(X = x) = \binom{n}{x} p^x (1 - p)^{n-x}, \quad x = 0, 1, \dots, n$$

We write $X \sim \text{Bin}(n, p)$

Tip: The number of successes in a series of n independent Bernoulli trials with success probability p has a $\text{Bin}(n, p)$ distribution

Theorem 6. Let $X \sim \text{Bin}(n, p)$. Then,

1. $\mathbb{E}(X) = np$
2. $\text{Var}(X) = np(1 - p)$

1.7.5 Binomial distribution in practice (cont.)

Many practical statistical situations can be treated as a sequence of coin flips. For example, we wish to conduct a survey of a large population to see what proportion of p is male, where p is unknown.

We can only know p if we observe everyone in the population, but we most likely can't do that. So we select n people at random from the population and record their gender. We assume that each person is chosen with equal probability.

If we allow the same person to be selected more than once, then the two situations are exactly the same. Consequently, if X is the total number of males in the group of n , then $X \sim \text{Bin}(n, p)$.

What if we never select the same person twice?

For a smaller population we should use a more complicated urn model to describe this experiment. (Select people without replacement without noting any order). But for a larger population, selecting a person twice is unlikely, so a Binomial model is a good model.

1.7.6 Conclusions!

- A Bernoulli trial is a random experiment with only two possible outcomes (success/failure)
- The total number of successes in n independent Bernoulli trials with success probability p has a $\text{Bin}(n, p)$ distribution. The expectation and variance are np and $np(1 - p)$ respectively.

1.8 Lecture 12 - Uniform and normal distribution

The simplest continuous distribution is the uniform distribution

Definition 10. A random variable X is said to have **uniform** distribution on the interval $[a, b]$ if its pdf is given by

$$f(x) = \frac{1}{b-a}, \quad a \leq x \leq b$$

(and $f(x) = 0$ otherwise). We write $X \sim U[a, b]$

1.8.1 Intepreteation

The random variable $X \sim U[a, b]$ can model a randomly chosen point from the interval $[a, b]$, where each choice is equally likely.

Theorem 7. Let $X \sim U[a, b]$. Then,

1. $\mathbb{E}(X) = \frac{a+b}{2}$
2. $\text{Var}(X) = \frac{(b-a)^2}{12}$

1.8.2 Normal Distribution

We introduce the most important distribution in the study statistics: the normal (or Gaussian) distribution.

Definition 11. A random variable X is said to have a **normal** distribution with parameters μ (expectation) and σ^2 (variance) if its pdf is given by

$$f(x) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{1}{2}\left(\frac{x-\mu}{\sigma}\right)^2}, \quad x \in \mathbb{R}$$

Theorem 8. Let $X \sim N(\mu, \sigma^2)$. Then

1. $\mathbb{E}(X) = \mu$
2. $\text{Var}(X) = \sigma^2$

if $\mu = 0$ and $\sigma = 1$, the distribution is called the **standard normal distribution**. Its pdf is denoted with φ

$$\varphi(x) = \frac{1}{\sqrt{2\pi}} e^{-\frac{x^2}{2}}$$

The respective cdf is denoted with capital phi Φ

1.8.3 Standardisation

Theorem 9. If Z has standard normal distribution, then $X = \mu + \sigma Z$ has a $N(\mu, \sigma^2)$ distribution. Consequently, if $X \sim N(\mu, \sigma^2)$ then the standardised random variable

$$Z = \frac{X - \mu}{\sigma}$$

has a standard normal distribution.

1.9 Lecture 13 - Multiple Random Variables

1.9.1 Experiments with more than one random variable

Most random experiments involve more than one random variable:

1. We select $n = 10$ people randomly and observe their heights, thus X_1, \dots, X_n can represent their individual heights.
2. We toss a coin repeatedly. Let $X_i = 1$, if the i th toss is Heads, $X_i = 0$ otherwise. The experiment describes a sequence of X_1, X_2, \dots of Bernoulli random variables.
3. We randomly select a person from a large population and measure mass (X) and height (Y).
4. We simulate 10,000 realisations from a standard normal distribution using `rnorm()` function—`rnorm()`— function. Let $X_1, \dots, X_{10,000}$ be the corresponding random variables.

1.9.2 Joint probability mass function

To describe the joint behaviour of a random variables X_1, \dots, X_n , we need to specify their **joint distribution**

Definition 12. The **joint pmf** of X_1, \dots, X_n (discrete) is the function f defined by

$$f(x_1, \dots, x_n) = \mathbb{P}(X_1 = x_1, \dots, X_n = x_n)$$

Which we sometimes write f_{x_1, \dots, x_n} instead of f to show that its a pmf of multiple random variables. And for sake of notation, we can denote $X = (X_1, \dots, X_n)$.

1.9.3 Calculating probabilities

We can calculate the probability that X lies in some set B via summation as

$$\mathbb{P}(X \in B) = \sum_{x \in B} f(x)$$

This is simply a consequence of the sum rule.

We can find the pmf of X_i by summing the join pdf over all possible values of the other variables

$$f_x(x) = \mathbb{P}(X = x) = \sum_y \mathbb{P}(X = x, Y = y) = \sum_y f_{x,y}(x, y)$$

1.9.4 Joint probability density function

For the continuous case, we need to replace joint pmf with joint **pdf**

Definition 13. The **joint pdf** f of continuous random variables X_1, \dots, X_n summarised as X is the positive function with total integral 1 such that

$$\mathbb{P}(X \in B) = \int_{x \in B} f(x) dx \text{ for all sets } B$$

Note that the integral in Definition 13 is a multiple integral and needs to be evaluated in n -dimensional volume ($|B| = n?$)

Figure 1.3: Left: a two-dimensional joint pdf of random variables X and Y . Right: the volume under the pdf corresponds to $\mathbb{P}(0 \leq X \leq 1, Y \geq 0)$

1.9.5 Independent discrete random variables

Two discrete random variables X and Y are said to be independent if the events $\{X = x\}$ and $\{Y = y\}$ are independent for every choice of x and y ; that is,

$$\mathbb{P}(X = x, Y = y) = \mathbb{P}(X = x)\mathbb{P}(Y = y)$$

Meaning that any information about the outcome of X does not provide extra information about Y . Formalizing:

Definition 14. Random variables X_1, \dots, X_n with joint pmf or pdf f are said to be **independent** if

$$f(x_1, \dots, x_n) = f_{x_1}(x_1) \cdots f_{x_n}(x_n)$$

$\forall x_1, \dots, x_n$ where $\{f_{x_i}\}$ are the marginal pdfs.

1.9.6 Bivariate Standard Normal Distribution

Suppose X and Y are independent and both have a standard normal distribution. We say that (X, Y) has a **bivariate standard normal distribution**. It's joint pdf is

$$f(x, y) = f_x(x)f_y(y) = \frac{1}{\sqrt{2\pi}}e^{-\frac{1}{2}x^2} \frac{1}{\sqrt{2\pi}}e^{-\frac{1}{2}y^2} = \frac{1}{\sqrt{2\pi}}e^{-\frac{1}{2}(x^2+y^2)}$$

Figure 1.4:

1.9.7 Simulating from the bivariate standard normal distribution

We can simulate copies $X_1, \dots, X_n \sim iid N(0, 1)$ and $Y_1, \dots, Y_n \sim iid N(0, 1)$ and plot the pairs $(X_1, Y_1), \dots, (X_n, Y_n)$ to understand the joint distribution. The following lines of R code produce the scatter plot of simulated data.

```
> x <- rnorm(2000)
> y <- rnorm(2000)
> plot(y~x, xlim = c(-3,3), ylim = c(-3,3))
```

Figure 1.5:

In the univariate case:

$$\mathbb{E}[h(X_1, \dots, X_n)] = \sum_{x_1, \dots, x_n} h(x_1, \dots, x_n) f(x_1, \dots, x_n)$$

Where the sum is taken over all possible values of (x_1, \dots, x_n) . In the continuous case replace the sum with a multiple integral.

Theorem 10. Let X_1, \dots, X_n be random variables with expectations, μ_1, \dots, μ_n . Then,

$$\mathbb{E}[a + b_1X_1 + b_2X_2 + \cdots + b_nX_n] = a + b_1\mu_1 + \cdots + b_n\mu_n$$

$\forall a, b_1, \dots, b_n$. Also, for independent random variables

$$\mathbb{E}(X_1X_2 \cdots X_n) = \mu_1\mu_2 \cdots \mu_n$$

1.9.8 Covariance

The covariance is a measure of the amount of linear dependency between two random variables.

Definition 15. The **covariance** of two random variables X and Y with expectations $\mathbb{E}X = \mu_X$ and $\mathbb{E}Y = \mu_Y$ is defined as

$$\text{Cov}(X, Y) = \mathbb{E}[X - \mu_X)(Y - \mu_Y)]$$

As easy reference for the properties of variance and covariance.

Theorem 11. For random variables X, Y and Z and constants a and b , we have

1. $Var(X) = \mathbb{E}(X^2) - (\mathbb{E}(X))^2$
2. $Var(a + bX) = b^2 Var(X)$
3. $Cov(X, Y) = \mathbb{E}(XY) - \mathbb{E}(X)\mathbb{E}(Y)$
4. $Cov(X, Y) = Cov(Y, X)$
5. $Cov(aX + bY, Z) = aCov(X, Z) + bCov(Y, Z)$
6. $Cov(X, X) = Var(X)$
7. $Var(X + Y) = Var(X) + Var(Y) + 2Cov(X, Y)$
8. If X and Y are independent, then $Cov(X, Y) = 0$

Combining properties (7) and (8) we find that if X and Y are independent, then the variance of the sum is equal to the sum of their variances.

Theorem 12. Let X_1, \dots, X_n be independent random variables with expectations μ_1, \dots, μ_n and variances $\sigma_1^2, \dots, \sigma_n^2$. Then,

$$Var(a + b_1X_1 + b_2X_2 + \dots + b_nX_n) = b_1^2\sigma_1^2 + \dots + b_n^2\sigma_n^2$$

$\forall a, b, \dots, b_n$

1.10 Lecture 14 - The Law of Large Numbers and the Central Limit Theorem

1.10.1 Law of Large Numbers

Theorem 13. The average of a large number of iid random variables tends to their expectation as the sample size goes to infinity.

If we want to say something about the expectation of a random variable. We can simulate lots of independent copies and take the average, to give a good approximation of the expectation.

Example 7. Let $U \sim U(0, 1)$. What is the expectation of \sqrt{U} ? We know that the expectation is $\frac{1}{2}$. Would the expectation of \sqrt{U} be $\sqrt{\frac{1}{2}}$? We can determine the expectation exactly, but we can generate a good approximation with the law of large numbers. We just simulate a lot of uniform numbers and take their square roots, then average them. In R this looks like

```
> u <- runif(10e6)
> x <- sqrt(u)
> mean(x)
```

Repeating this gives 0.666... consistently. The true expectancy is $\frac{2}{3} < \sqrt{\frac{1}{2}}$.

1.10.2 Central Limit Theorem

Theorem 14. The sum of a large number of iid random variables approximately has a normal distribution. $\forall X_1, X_2, \dots$ sequences of iid random variables with finite expectation μ and finite variance σ^2 , the sum

$$S_n = X_1 + X_2 + \dots + X_n$$

has approximately a $N(n\mu, n\sigma^2)$ distribution.

1.10.3 Demonstration of the CLT via simulation

This is a truly remarkable result and is one of the greatest milestones in mathematics.

We can demonstrate using simulation.

Let X_1 be a $U[0, 1]$ random variable. The pdf is constant on the interval $[0, 1]$ and 0 elsewhere. If we simulate many independent copies of X_1 and take the histogram, the result will resemble the shape of the pdf (which is a consequence of the law of large numbers).

What about the pdf of $S_2 = X_1 + X_2$? We can generate many copies of both X_1 and X_2 and add them up to create a histogram. In R this would look

```
> x1 <- runif(10e6)
> x2 <- runif(10e6)
> hist(x1 + x2, breaks=100, prob=T)
```

Figure 1.6: Histogram for the sum of 2 independent uniform random variables

We can do the same thing for sums of 3 and 4 uniform numbers

Figure 1.7: The histograms for the sums of 3 (left) and 4 (right) uniforms are closely in agreement with normal pdfs

The CLT suggests that **taking linear combinations of independent normal random variables, gives again a normal random variable**. This is true. Observe.

Theorem 15. Let X_1, X_2, \dots, X_n be independent normal random variables with expectations μ_1, \dots, μ_n and variances $\sigma_1^2, \dots, \sigma_n^2$. Then for any numbers a, b_1, \dots, b_n the random variable

$$Y = a + b_1 X_1 + b_2 X_2 + \dots + b_n X_n$$

has a normal distribution with expectation $a + \sum_{i=1}^n b_i \mu_i$ and variance $\sum_{i=1}^n b_i^2 \sigma_i^2$

1.10.4 Conclusion

- Multiple random variables are specified by their joint cdf or pmf/pdf
- For independent random variables the joint pmf/pdf is the product of the marginal ones.
- The expectation of the sum of random variables is equal to the sum of their expectations.
- The variance of the sum of independent random variables is equal to the sum of their variances.
- The expectation of the product of independent random variables is equal to the product of their expectations.
- The law of large numbers says that the average of a large number of iid random variables is approximately equal to their expectation.
- The central limit theorem says that the sum of large number of iid random variables is approximately normal.
- Any linear combination of independent normal random variables is again normal.

Chapter 2

Ethics

Working backwards.

2.0.1 Animal Experimentation

A possible way to avoiding the moral issues of experimenting on humans is to experiment on animals instead. The argument of animals not being "people" is not very persuasive, similar arguments have been used (they are not worthy of moral consideration).

What constitutes ethical experimentation on animals however? Respect for animals autonomy isn't wildly considered as an ethical principle, but *consequential ethical reasoning* does not give researches the ability to do whatever they want.

2.0.2 The Efficacy Argument

The efficacy argument fundamentally goes like this:

If animals are not a close physiological match to humans then there seems to be little to no point of using them as reference for experimentation. On the other hand, if they are good representation of humans; then its unethical to experiment on them. There are *prima facie*, good reasons not to inflict pain upon them (without a very good reason)

2.0.3 Animal Experimentation (cont.)

A pure utilitarian or consequentialist reasoning is the best way we have for framing ethical discussions of animal experimentation. The areas in which animals are used as research subjects can be thought of as falling into 3 **broad** areas

- Industrial Research
- "Pure" scientific research
- Biomedical research

A common idealized view of science is of objective research. Carrying out objective research, presenting in an objective free way. Ignoring other influences. There are merits to this, but there also various problems.

2.0.4 Experiment Duplications

Why are so many experiments not duplicated?

1. Because there is little motivation for researchers to replicate other results rather than finding something new.
2. Experiments cost money.
3. In the case for drug experiments, because of the intellectual property system. A researched needs to gain permission of the owner of the patent and is commonly withheld or given with significant restrictions.

2.0.5 Ethical Failures

Reasons for Ethical failures

1. Desire for grant money
2. Desire for personal money
3. Pressure to publish
4. Avoidance of hard work
5. Already "knowing" the right answer
6. Desire for fame
7. Political/Religious pressure

2.0.6 Retin-A-Case

2.1 Lecture Don't Know

2.1.1 Most important suggestions

- **Give yourself time to do lots of redrafting.** The more time you take to redraft the better your essay will be.
- You can't proofread your work until 24 hrs after you wrote something
- Going back to your essay again and again, rewriting something more simply and directly is essential to a good essay
- Each time you read over it as if you are reading it for the first time as a very critical reader. Edit all parts that are unclear or can be expressed more succinctly and logically.
- You have the opportunity for tutors to comment on the parts you are unsure of in the tutorial the week after the Wakefield Paper (last week before midsem break).
- You can only take that opportunity if you have a draft ready before your tutorial.

2.1.2 Retractions and Fraud in Science Publications

From the lecture slides:

"A survey of the 2046 articles in the PubMed database that had been marked as retracted by 3 May, 2010 found that two-thirds of retracted life-science papers were stricken from the scientific record because of misconduct such as fraud or suspected fraud. Other types of misconduct - duplicate publication and plagiarism - accounted for 14% and 10% of retractions, respectively. Only 21% of the papers were retracted because of error."

2.1.3 Two most common ethical issues facing scientists